

# RFC: Compact Telemetry Protocol (CTP)

Version 1.0

**Prepared By:**

Ahmed Mohamed Al Amin (22P0137)  
Mohamed Yehia Zakaria (22P0064)  
Jana Hany Elshafie (22P0235)  
Jana Sameh Samir (22P0237)  
Seif Elhusseiny (22P0215)  
Youssef Tarek Kamal (22P0236)

**Delivered To:**

Dr. Ayman Bahaa Eldin  
Dr. Karim Emara

Faculty of Engineering, Ain Shams University  
Computer Engineering and Software Systems Department  
CSE361 — Computer Networking

Fall 2025

<b>Transport Protocol</b>	UDP
<b>Default Port</b>	12000
<b>Byte Order</b>	Big Endian (Network)
<b>Status</b>	Experimental

## 1 Introduction

The **Compact Telemetry Protocol (CTP)** is a lightweight, binary application-layer protocol designed for constrained IoT sensors operating over UDP. To address bandwidth limitations and battery constraints, CTP eliminates traditional TCP connection overhead, using a minimal 11-byte bit-packed header to maximize efficiency and support high-frequency telemetry. The protocol adopts a strict “fire-and-forget” model suitable for idempotent sensor data (e.g., temperature, humidity, voltage). It does not support retransmission (ARQ); instead, the Server detects gaps in sequence numbers without attempting recovery.

Despite its simplicity, CTP preserves observability and data integrity through the following architectural constraints:

- **Latency Handling:** UDP jitter is managed entirely at the application layer using a timestamp-driven Server Reordering Buffer, preventing head-of-line blocking and ensuring in-order processing.
- **Data Representation:** All multi-byte fields **MUST** be encoded in Network Byte Order (Big Endian). Timestamps use 32-bit integers relative to a Custom Epoch (Dec 1, 2025) to preserve millisecond-level precision while minimizing header size.
- **Integrity:** Every packet includes a mandatory 16-bit checksum, enabling the Server to detect transmission corruption and discard invalid data.

## 2 Protocol Architecture

### 2.1 Entities & Operational Modes

CTP follows a unidirectional **Push Model** in which the Sensor (Client) continuously transmits telemetry to a Collector (Server). The Client operates in one of two modes:

1. **Direct Mode:** The Client sends a packet immediately after each sensor reading to minimize latency.
2. **Batching Mode:** The Client aggregates  $N$  readings into a local buffer and transmits them as a single payload once the buffer is full ( $N \leq 16$ ), maximizing throughput.

### 2.2 Reliability & Liveness

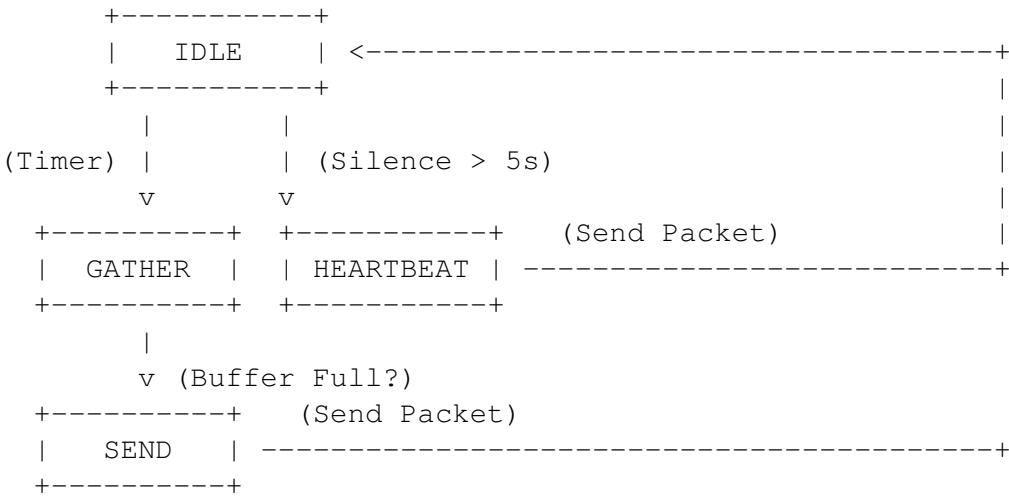
**Liveness Monitoring:** If the Client generates no data for `HEARTBEAT_INTERVAL` (5 s), it sends a `HEARTBEAT` packet. The Server marks a device `OFFLINE` if no packets are received for `LIVENESS_TIMEOUT` (10 s).

**Jitter Handling:** The Server employs a **Reordering Buffer**. Packets are queued upon arrival and sorted by timestamp rather than processed immediately. The buffer is flushed to application logic when its size exceeds `FLUSH_THRESHOLD` (20 packets), ensuring correct temporal ordering under UDP jitter.

## 2.3 Finite State Machines (FSM)

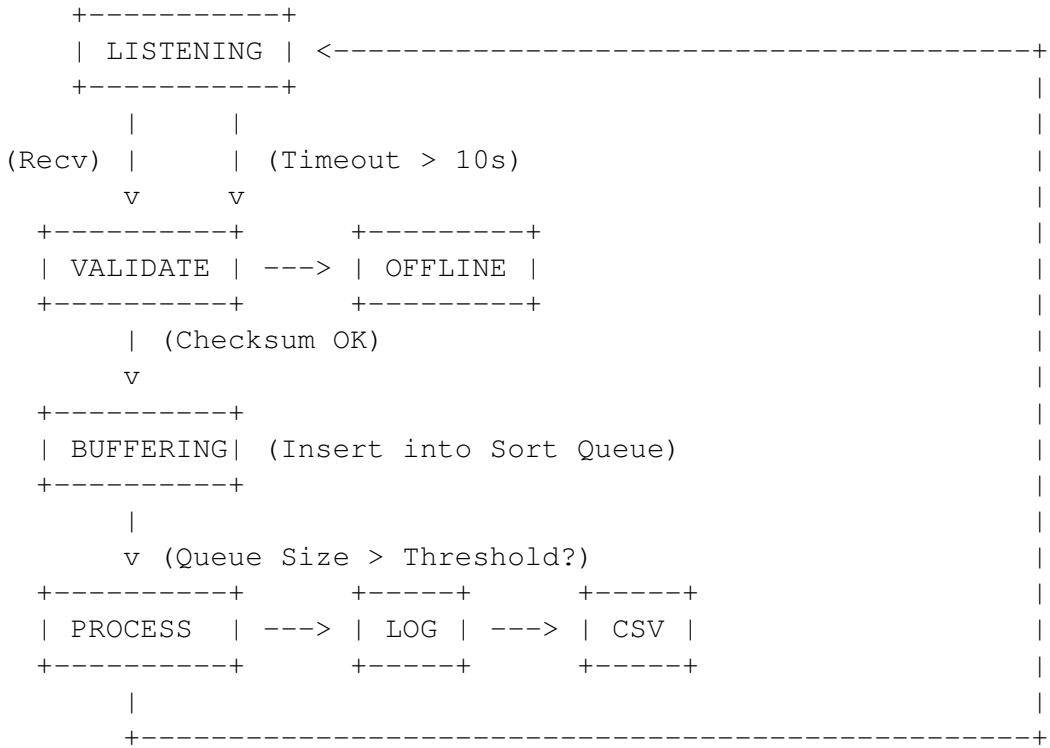
### Client FSM

The Client alternates between reading sensor data and managing network transmission.



### Server FSM

The Server manages packet validation, jitter buffering, and device state tracking.



## 2.4 Sequence Diagram (Normal Flow)

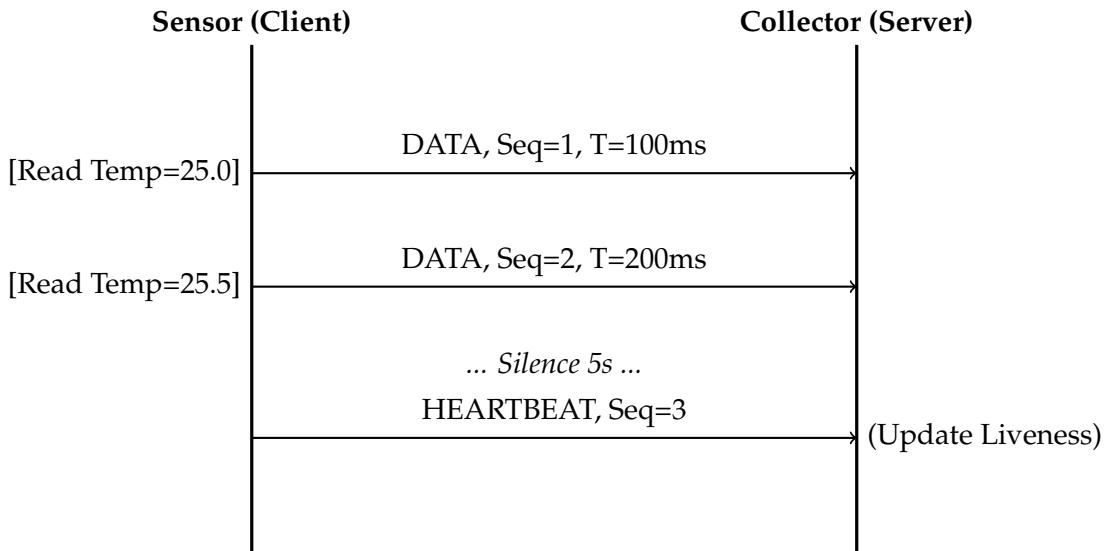


Figure 1: Typical communication flow showing data transmission and heartbeat.

## 3 Message Formats

### 3.1 General Encoding Rules

- **Byte Order:** All multi-byte integers MUST be transmitted in Network Byte Order (Big-Endian). In Python struct syntax, this is denoted by the ! prefix.
- **Alignment:** Fields are packed tightly without padding bytes.
- **Total Header Overhead:** 11 Bytes (9 Byte Fixed Header + 2 Byte Checksum). This strictly satisfies the requirement of  $\leq 12$  bytes.

### 3.2 Fixed Header Structure

Every packet (whether DATA or HEARTBEAT) begins with the following 11-byte sequence.

Offset	Field Name	Size (Bits)	Type	Description
0	Device ID	16	uint16	Unique identifier for the sensor node.
2	Seq Num	16	uint16	Monotonic counter. Wraps at 65,535.
4	Timestamp	32	uint32	Milliseconds since Custom Epoch (see 3.3).
8	Ver / Type	8	uint8	Bit-Packed: [Ver: 4 bits   Type: 4 bits].
9	Checksum	16	uint16	Additive sum of (Header + Payload).

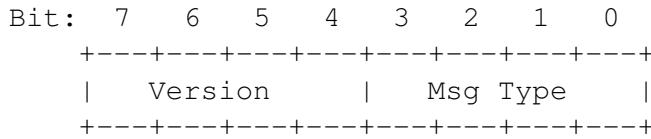
**Python Struct Format:** !H H I B H

*Note: The first four fields are packed via !HHIB, and the checksum !H is appended after being calculated over the specific payload.*

### 3.3 Field Definitions & Handling Logic

#### 3.3.1 The "Ver / Type" Packed Byte (Offset 8)

To optimize space, the Protocol Version and Message Type are combined into a single byte using bitwise operations.



- **Encoding Logic:** `packed_byte = (Version << 4) | (MsgType & 0x0F)`
- **Decoding Logic:**
  - `Version = (packed_byte >> 4) & 0x0F`
  - `MsgType = packed_byte & 0x0F`

#### Supported Message Types:

- `0x01` = **DATA** (Contains Sensor Readings)
- `0x02` = **HEARTBEAT** (Empty Payload, Keep-alive only)

#### 3.3.2 Timestamp Compression (Offset 4)

Standard Unix timestamps (seconds since 1970) do not provide enough precision for jitter analysis, but sending full 64-bit milliseconds is wasteful.

##### Optimization Strategy:

- **Custom Epoch:** We define a custom epoch `TIMESTAMP_OFFSET = 1764547200` (Corresponding to Dec 1, 2025).
- **Truncation:** The timestamp is calculated as `(CurrentTime - Offset) * 1000`.
- **Wrapping:** The result is masked to 32 bits (`& 0xFFFFFFFF`).

This allows us to track millisecond-level precision within a 4-byte field, which is critical for the `latency_ms` and `jitter_ms` calculations in the Server.

#### 3.3.3 Checksum (Offset 9)

The 2-byte checksum provides basic integrity verification.

- **Calculation:** `sum(byte_array) & 0xFFFF`
- **Scope:** Calculated over the Initial Header (9 bytes) + The Variable Payload.
- **Validation:** The server re-calculates the sum of received bytes (excluding the checksum field itself) and compares it to the received checksum value. Mismatches result in immediate packet drops.

### 3.4 Payload Formats

#### 3.4.1 Heartbeat Payload

- **MsgType:** `0x02`
- **Length:** 0 Bytes
- **Content:** None. The header timestamp serves as the “last seen” proof.

### 3.4.2 Data Payload (Variable)

To support the optional batching requirement universally, all DATA payloads use a Count-Prefixed structure. Even a single reading is treated as a batch of size 1.

- **Structure:**
  - **Byte 0:** Count (uint 8) - Number of readings ( $N$ ) in this packet.
  - **Byte 1..M:** Array of Readings ( $N \times 12$  bytes).
- **Reading Format (12 Bytes per Reading):** ! f f f (Temperature, Humidity, Voltage).
- **Constraint:** The total payload size MUST NOT exceed 200 Bytes.

$$1 + (N \times 12) \leq 200 \implies N_{\max} = 16 \text{ readings} \quad (1)$$

## 4 Communication Procedures

### 4.1 Session Lifecycle (Stateless)

**Initiation:** CTP operates using a strictly stateless “fire-and-forget” model with no connection handshake (no SYN/ACK). The Client begins transmitting immediately upon startup. When the Server receives the first valid packet from a new DeviceID, it allocates the necessary session state, including a jitter buffer and sequence tracking structures.

**Termination:** There is no explicit teardown (no FIN). If the Server receives no packets from a device for more than LIVENESS\_TIMEOUT (10.0 s), that device is marked OFFLINE. Any remaining packets in the jitter buffer are immediately flushed and processed to avoid data loss.

### 4.2 Data Exchange

**Transmission:** In Direct Mode, Clients transmit readings at the configured REPORTING\_INTERVAL (default: 1.0 s). In Batching Mode, readings are accumulated locally and transmitted as a single aggregate payload once the buffer size reaches BATCH\_SIZE.

**Keep-Alive:** If the Client has no data to send for more than HEARTBEAT\_INTERVAL (5.0 s), it sends a HEARTBEAT packet to prevent the Server from marking the device as inactive.

### 4.3 Error Handling & Recovery

**Corruption:** The Server verifies each packet’s 16-bit checksum. Any checksum mismatch results in the packet being silently dropped.

**Packet Loss:** CTP does not support retransmission (ARQ). Missing sequence numbers are detected by the Server and logged as *Gaps*, providing visibility into packet loss without attempting recovery.

**Duplicate Suppression:** The Server maintains a history of the last 500 processed sequence numbers. If a packet arrives with a sequence number that already exists in this history, it is classified as a duplicate and its payload is ignored.

## 5 Reliability & Performance Features

### 5.1 Loss Tolerance Strategy

CTP follows a **loss-tolerant by design** philosophy. It intentionally does *not* implement Automatic Repeat Request (ARQ) or retransmission mechanisms. **Rationale:** In real-time telemetry

systems, fresh data is more valuable than delayed reliability. Retransmitting stale sensor values wastes bandwidth that should instead carry current readings.

**Gap Reporting:** The Server detects packet loss by monitoring discontinuities in the sequence number:

$$Seq_{\text{current}} - Seq_{\text{last}} > 1$$

Any such gap is logged for later network-health diagnostics.

## 5.2 Jitter Compensation (Windowing)

Because UDP provides no ordering guarantees, the Server implements a **Reordering Window**.

**Mechanism:** Incoming packets are not processed immediately. They are inserted into a Sort Queue and ordered by their timestamp.

**Flush Threshold:** To avoid unbounded buffering, the queue is flushed only when:

- its size exceeds FLUSH\_THRESHOLD (default: 20 packets), or
- the session terminates.

This ensures the application layer receives events in correct temporal order even when packets arrive out of order.

## 5.3 Duplicate Suppression

To prevent double-counting caused by duplicated UDP deliveries or link-layer retries, the Server maintains a sliding history of the last 500 processed sequence numbers.

**Rejection Logic:** If an incoming SeqNum already exists in this history, the packet is classified as a **Duplicate** and its payload is discarded.

## 5.4 Timeout Selection

Timer values were chosen to balance bandwidth efficiency, responsiveness, and NAT stability:

- **Heartbeat Interval (5.0 s):** Low transmission overhead (< 1% duty cycle), while keeping NAT bindings alive.
- **Liveness Timeout (10.0 s):** Defined as:

$$2 \times \text{HEARTBEAT\_INTERVAL}$$

This allows the Server to tolerate a single lost heartbeat without incorrectly marking the device as OFFLINE.

# 6 Experimental Evaluation Plan

## 6.1 Measurement Methodology

Performance evaluation is conducted using an automated Linux-based test suite (`Run.sh`).

**Network Emulation:** The Linux `netem` module is used to inject controlled latency, jitter, loss, and duplication on the `lo` interface.

**Verification:** Each experiment captures a `.pcap` trace using `tcpdump` to validate header correctness and on-wire behavior.

**Timing:** Server-side processing time is measured using `time.perf_counter()` to record per-packet computational cost with microsecond precision.

## 6.2 Key Metrics

For each received packet, the Server logs the following metrics to a CSV file:

- **Latency (ms):**

$$T_{\text{arrival}} - T_{\text{sent}}$$

based on synchronized clocks or relative timestamps.

- **Jitter (ms):**

$$|Latency_{\text{current}} - Latency_{\text{prev}}|$$

- **Sequence Gaps:** Missing packets inferred from:

$$Seq_{\text{current}} - Seq_{\text{last}} - 1$$

- **Duplicate Rate:** Percentage of packets rejected by the Duplicate Suppression logic.
- **CPU Cost:** Time required to parse, validate, reorder, and log a packet.

## 6.3 Test Scenarios & Netem Configurations

The following scenarios are executed automatically by Run.sh to evaluate key protocol behaviors:

Scenario	Netem Command	Objective
Baseline	none	Establish reference metrics for latency and throughput (1 s reporting interval).
Packet Loss	loss 5%	Verify that the Server detects Gaps without interrupting operation.
Heavy Loss	loss 30%	Stress-test gap detection logic under severe degradation.
Jitter	delay 100ms 10ms	Validate that the Reordering Buffer correctly restores timestamp order.
Duplication	duplicate 20%	Confirm that Duplicate Suppression filters redundant packets.
Batching	none	Evaluate throughput efficiency when sending packets containing 5 aggregated readings.

## 7 Example Use Case Walkthrough

### 7.1 Narrative Description

**Startup ( $T = 0$ ):** Device 101 powers on and initializes its sequence counter to 1.

**Data Transmission ( $T = 1000 \text{ ms}$ ):** The sensor reads  $25.5^\circ\text{C}$  and sends a DATA packet with  $Seq = 1$ .

**Packet Loss ( $T = 2000 \text{ ms}$ ):** The sensor reads  $25.6^\circ\text{C}$  and sends  $Seq = 2$ , but the packet is dropped due to simulated network loss (5%).

**Gap Detection ( $T = 3000 \text{ ms}$ ):** The sensor transmits  $Seq = 3$ . The Server receives it, detects that the previous packet was  $Seq = 1$ , and logs a Gap of 1 packet.

**Idle & Heartbeat ( $T = 8000 \text{ ms}$ ):** The sensor becomes idle. After 5 seconds of silence, it sends a HEARTBEAT packet with  $Seq = 4$  to maintain liveness.

## 7.2 Trace Log (Server CSV Output)

Example CSV entries corresponding to the scenario:

Device ID	Seq	Msg Type	Gap?	Latency (ms)	Notes
101	1	DATA	0	2.05	Initial Packet. Session Created.
101	3	DATA	1	1.98	Gap Detected (Seq 2 lost).
101	4	HEARTBEAT	0	1.50	Keep-alive received.

## 8 Limitations & Future Work

### 8.1 Current Limitations

**Security:** The protocol lacks authentication and encryption. Device IDs and source IPs can be spoofed, enabling injection of false telemetry.

**Scalability:** The Server uses a Python dictionary for state tracking. While adequate for a small lab setup (4 clients), it may not scale to thousands of devices due to memory and locking constraints.

**Congestion Control:** Clients transmit at a fixed REPORTING\_INTERVAL and do not adjust rate in response to congestion. This risks contributing to network overload.

**Clock Synchronization:** Latency measurements assume synchronized clocks. Without NTP or PTP, one-way delay calculations in real-world deployments may be inaccurate.

### 8.2 Future Work

**Security Layer:** Add HMAC-SHA256 authentication to optional header fields to verify packet origin and prevent spoofing.

**Adaptive Rate Limiting:** Enhance the Client to dynamically adjust its REPORTING\_INTERVAL based on observed packet loss (via a new ACK message type).

**Dynamic Buffering:** Replace the fixed-size Reordering Buffer (20 packets) with a dynamically sized jitter buffer that scales with measured arrival-time variance.

## 9 References

1. Tanenbaum, A. S., & Wetherall, D. *Computer Networks*, 5th or 6th Edition, Pearson, 2011/2021.
2. Stallings, W. *Data and Computer Communications*, 10th Edition, Pearson, 2014.
3. Kurose, J. F., & Ross, K. W. *Computer Networking: A Top-Down Approach*, 8th Edition, Pearson, 2021.
4. Postel, J. "User Datagram Protocol," RFC 768, August 1980.
5. IEEE Computer Society. *IEEE Standard for Floating-Point Arithmetic*, IEEE 754-2019.